

High Level Design

Coding

- ① Problem Statement
- ② Functional Requirements
- ③ Non-functional Requirements
- ④ Scale Estimation
- ⑤ Architecture Design

SDE - 2 on High Positions - architectural design decisions

Any tier - 1 company - have at least 1 System Design round (HLD)

Staff Engineer @ Google.] 1 - 2 Crore per annum
10+ yrs in India

Interview Question for Staff Engineer @ Google → HLD round

List of strings → sort them in dictionary order

Cat
 dog
 apple
 High level design
 Scales

$\xrightarrow{\text{sort}}$

apple
 cat
 dog
 High level design,
 scales

Catch*

50 Petabytes of data

50,000,000 GB of data

v. Large Scale

Ryt - 8 bits

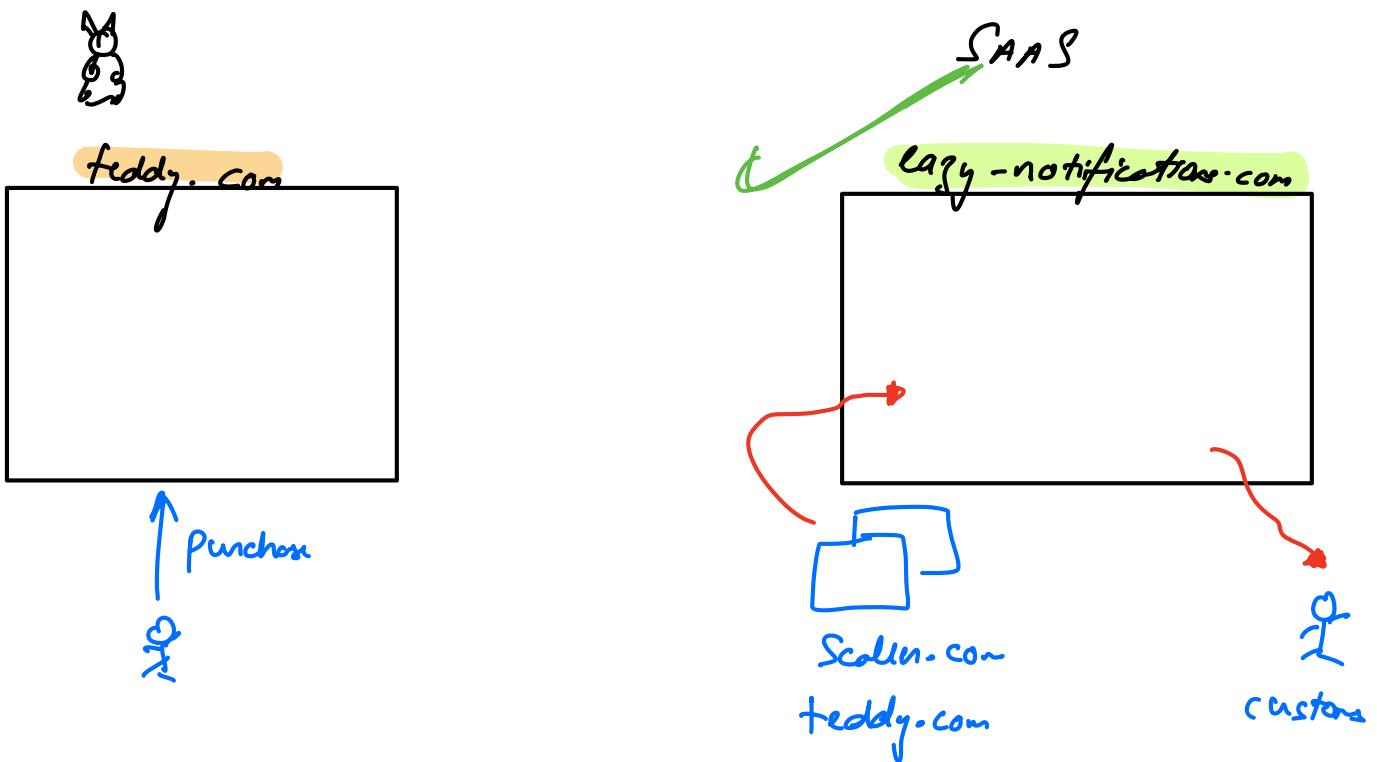
KB	- 1000 B
MB	- million B
GB	- billion
TB	- trillion
PB	- quintillion

Simple problems $\xrightarrow[\text{Planet Scale}]{\text{v. Large Scale}}$ v. difficult

High Level Design

① Problem Statement

Notification delivery service



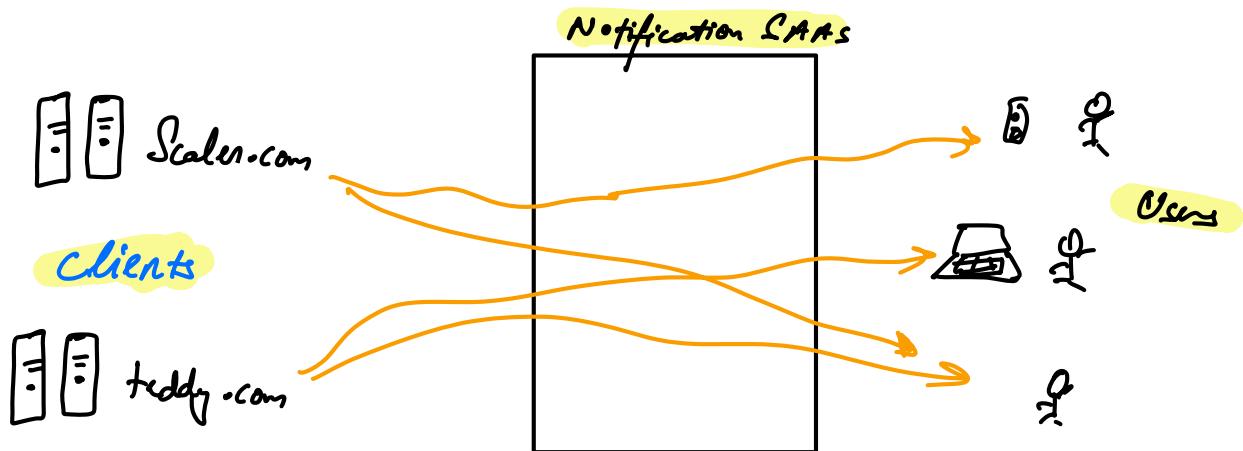
- Needs to send notifications to its customers
- Primary function - sell product
- Primary function is notifications.

② Functional Requirements

MVP features

Minimal Viable Product

- don't think of fancy features



Client Requirements

① Client should be able to send notifications to users

notify (_{sender} client-id, _{recipient} user-id, _{payload} message) → success / failure

② Support multiple delivery channels

SMS/Email/in-app

slack/whatsapp/Pigeone

Extensible / Pluggable

We should be able to add new delivery channels without changing infrastructure

- ③ client should be able to broadcast messages to a group of users

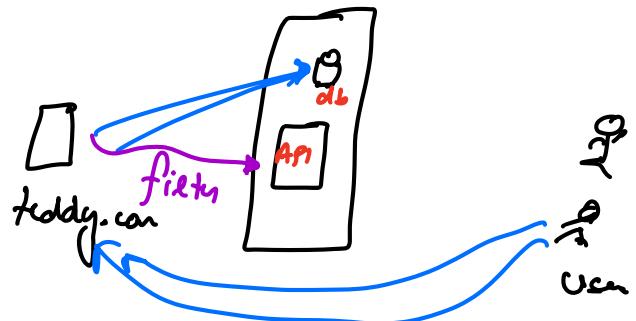
~~notify (client-id, [list of user-ids], message) → success / failure~~

bad idea!!

sender recipients payload

req Payload is too large

- Clients already store their customer contact info in our DB
- Clients can specify a "broadcast filter" to send bulk messages



broadcast (client-id, filter, message) → success / failure

↓
Example (User is subscribed to (1224)
or User.last-login > 24 hours)

User Requirements

- ① User should be able to receive notifications
 - ↳ in real time (if they are online)
 - ↳ on re-connecting (once they come back online)
- ② users should be able to customize their notification preferences

Select channels

- Email
- SMS
- WhatsApp
- Pigeons

Notification Type

- Marketing
- Trending / News
- Financial
- OTP

Hidden Requirements

① Rate limiting

Client Limiting

free : 100 notifications / day

Paid : 50,000 notifications / day

Enterprise : up to 50 M notifications / day

User Limiting

a user should not get more than — notifications in — time

② Observability

back in 2020 Google had
≈ 11 million servers

detecting that something is wrong

Metrics like

notifications sent / hour

of errors / minute

any CPU load

diagnose

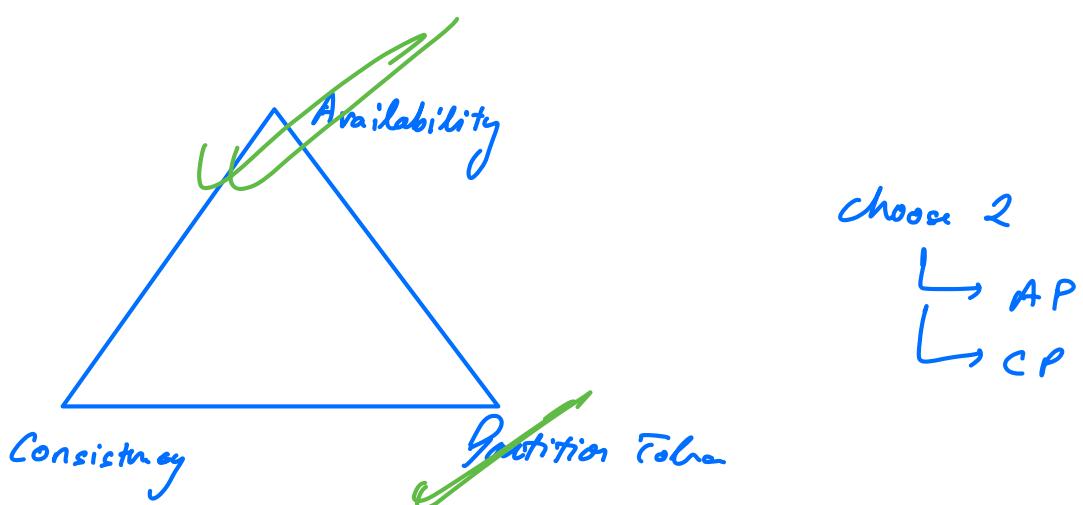
Search logs

verify the fix

③ Non-functional Requirements (Design Goals)

CAP / PACELC theorem

Eventual Consistency vs Immediacy Consistency
for some time your data can be stale



Imagine Mala has purchased a Teddy
the teddy is out for delivery

teddy.com → Mala
in worst case when systems are down
for some time when mala checks her phone she
gets no notification

but after some time Mala gets notified

usually

Mala gets notified immediately



- Compliance TRAI /FDR/...
- Priority OTPs should be prioritized over marketing notifications
- Exactly once delivery
 - ↳ 0 delivery ✗ *leaky*
 - ↳ the same notification just delivered 5 times ✗ *this should never happen*
 - exactly once ✓
- ~~Security~~ *discuss with security experts*

(4)

Scale Estimation

- ↳ it is okay to be a little off
- ↳ it is okay to overestimate

Total # users \approx 2 billion

Avg # of notifications / user / day \approx 10

$$\text{Avg } \# \text{ notifications / second} = 200,000$$

$$2 \times 10^9 \text{ users} * 10 \text{ notifications / user / day}$$

$$= 2 \times 10^{10} \text{ notifications / day}$$

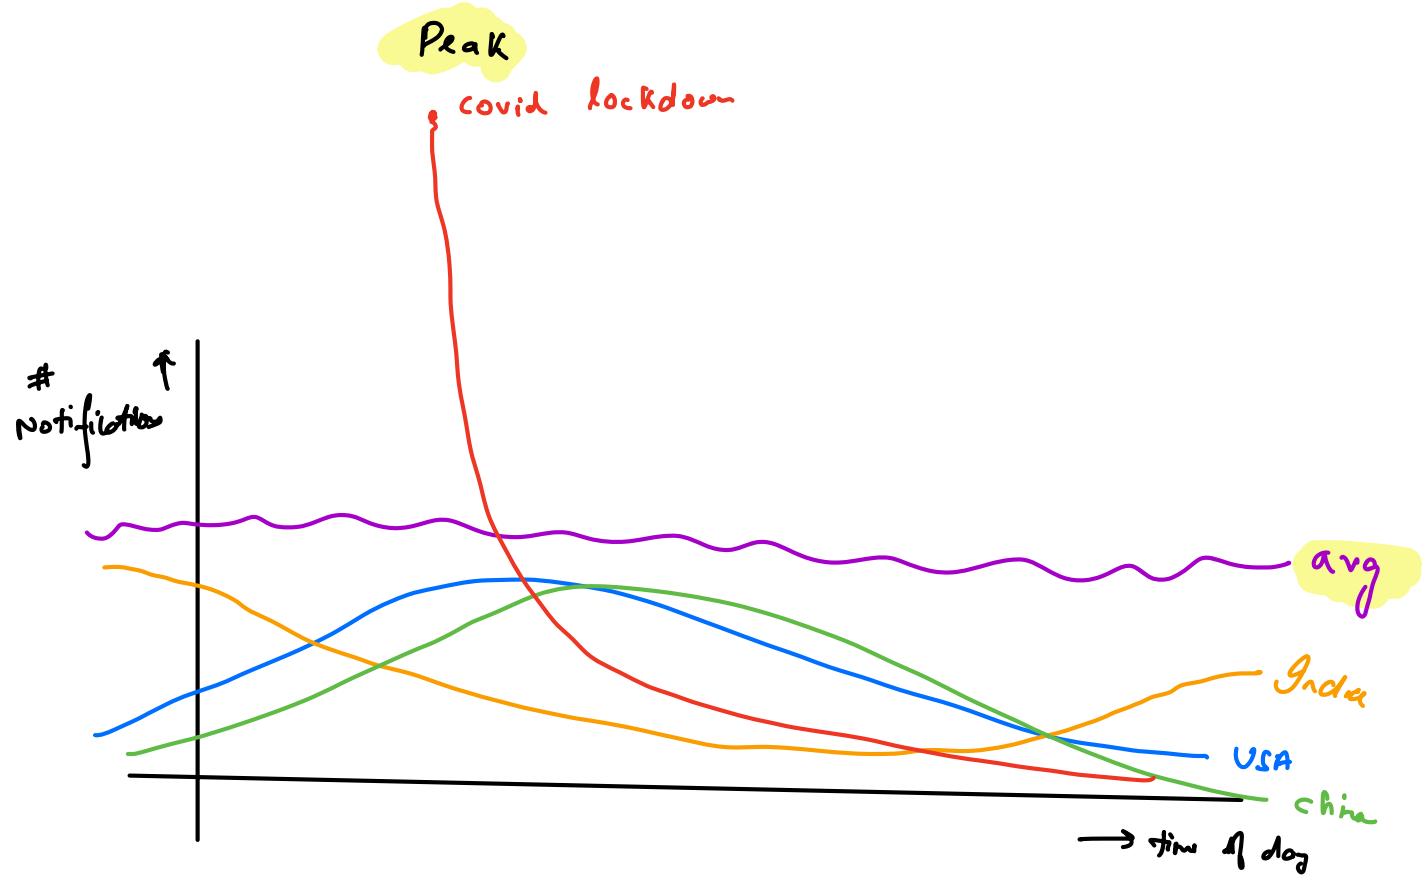
$$\hookrightarrow 60 * 60 * 24 \text{ seconds}$$

$$= 86400 \text{ seconds}$$

$$\hookrightarrow 10^5 \text{ seconds}$$

$$= \frac{2 \times 10^{10} \text{ notifications}}{10^5 \text{ seconds}}$$

$$\hookrightarrow 200,000 \text{ notifications / sec}$$



Peak could be 50x higher than avg

$$\begin{aligned}
 \text{Peak # notification /second} &= 200,000 \text{ avg/sec} \times 50 \\
 &= 10 \text{ million notifications /sec}
 \end{aligned}$$

Jargon

```
graph LR; Jargon --> LB[Load Balancer]; Jargon --> C[Caching]; Jargon --> NS[NoSQL]; Jargon --> Ellipsis[...]
```

Load Balancer
Caching
NoSQL
...

9:07 → 9:15

Scalor HLD Curriculum

- How backend systems work

DNS / vertical vs horizontal scaling / load balancing / sharding

- Load Balancing & Consistent Hashing

Healthcheck & Heartbeat / Role own /
DNS as load balancer / Untrusted Layer /
Consistent hashing - duplicate

- Caching

Browsn / CDN / Backend

Local vs Global / Single vs Distributed
Redis

Cache Invalidation

- Write Through
- TTL
- Write Back
- Write Around

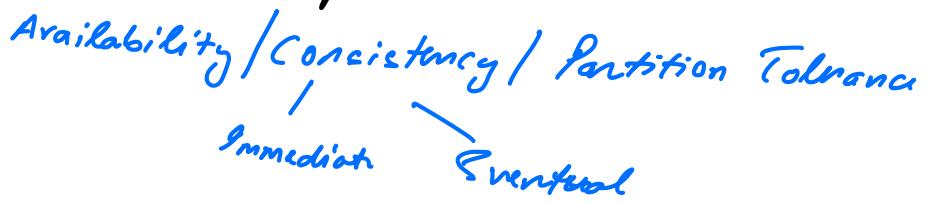
Cache Eviction

- FIFO
- LRU
- LFU
- MRU

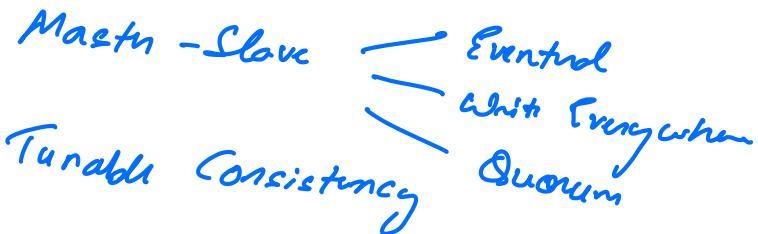
- Caching Case Studies

Scalor Code Judge / Scalor Leaderboard /
Facebook News Feed

• CAP theorem & Replication



CAP theorem / PACELC



• SQL vs NoSQL

Pros & Cons

Sharding Key



• Database Internals

Orchestrator / Config Management / Resend frames /
Shard Creation / Multi-master

• Database Internals

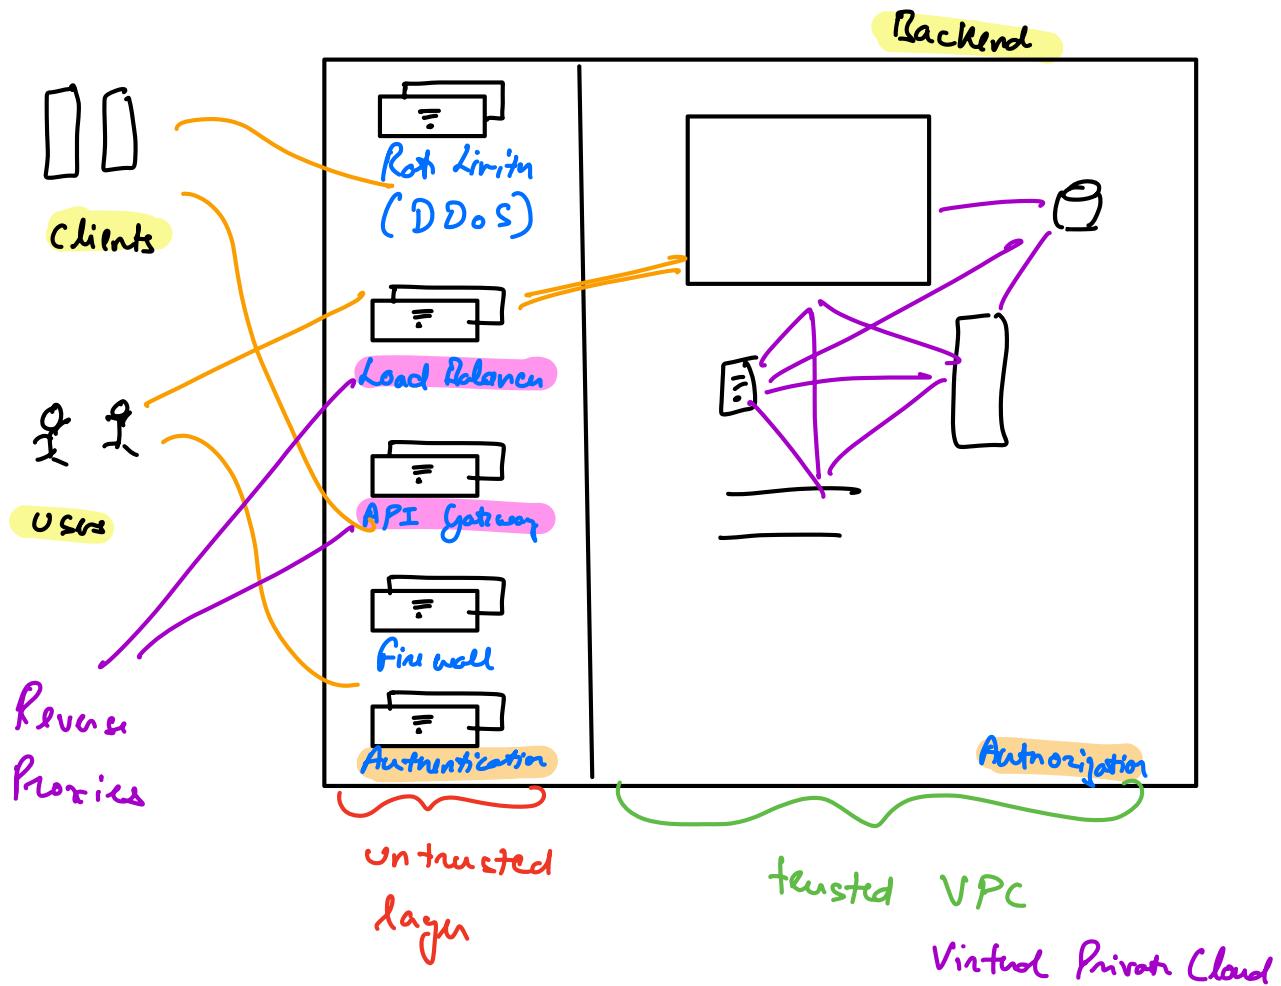
LSM Tree / Bloom Filters / Sparse Index /
Compaction / MemTable / WAL

- Case Study - Google Typahead
trie / Redis / Recency / Sampling / Ratcheting / Percolation
- Case Study - Messaging apps
Slack groups / WhatsApp / Facebook messenger
- Case Study - Kafka + Zookeeper
event driven architecture / Pub-Sub / topics & partitions / msg order
- Case Study - Elastic Search
Inverted Index / TF-IDF / Sharding / Tunable Latency
- Case Study - File Storage
upload / download / chunking
- Case Study - Uber
Quad Tree / traffic patterns / nearest driver
- Case Study - Rate Limiting
DDoS / Token bucket / Leaky Bucket / Sliding window

- Case Study - Id generation
idempotency / Unique / incremental
- Case Study - Video Streaming
Netflix / YouTube / Riotbox
Concurrency
ARoS / chunking / CDNs
- Microservice
 - Monolith vs Microservices vs Service Oriented
 - Breaking a Monolith
 - Observability / Distributed Tracing / Logging
 - Thundering herd / Cascading failure / Circuit Breaker
 - Distributed Transactions
 - Two Phase Commit / Saga / CQRS /
 - Orchestration & Chorography
 - Event Sourcing

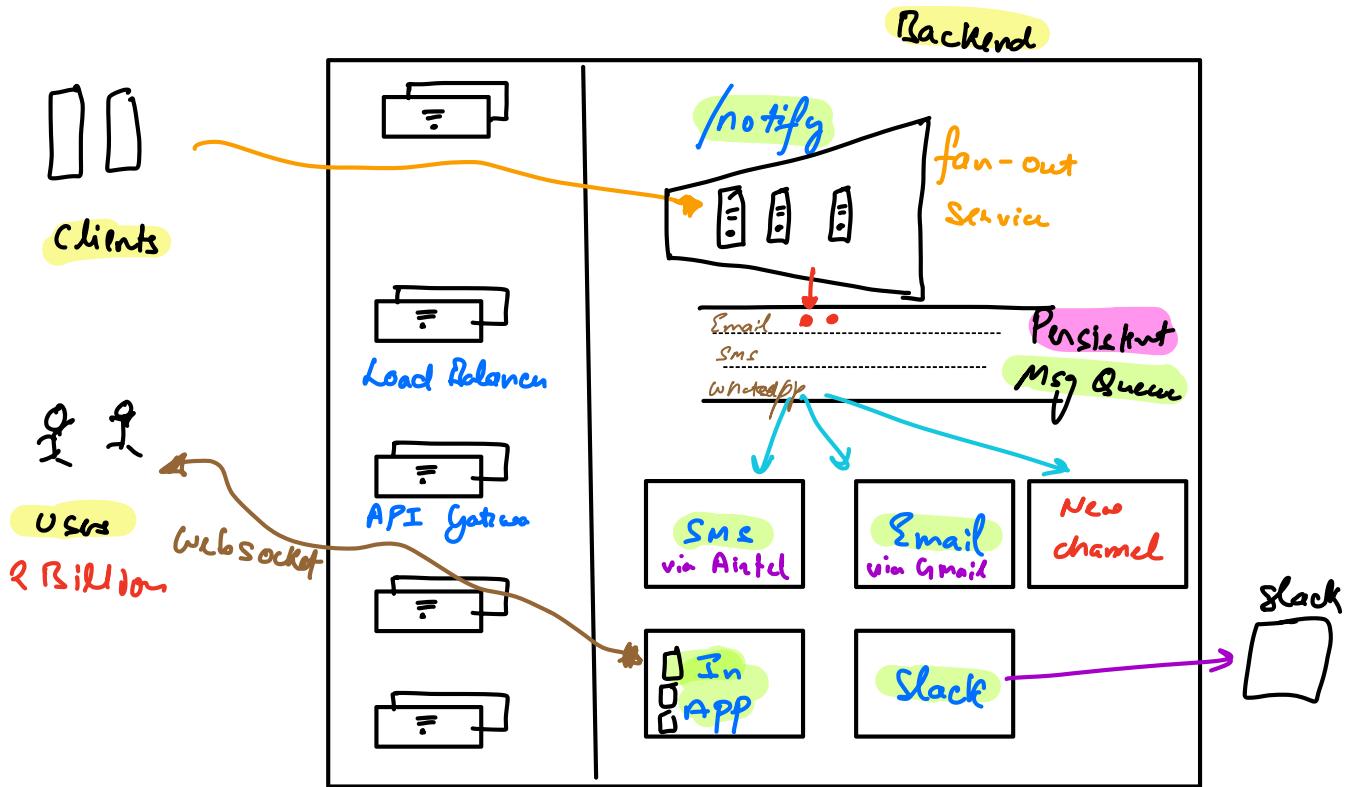
⑤ Architecture Design

Generic Design



①

Clients should be able to send notifications
 User should be able to receive notifications.

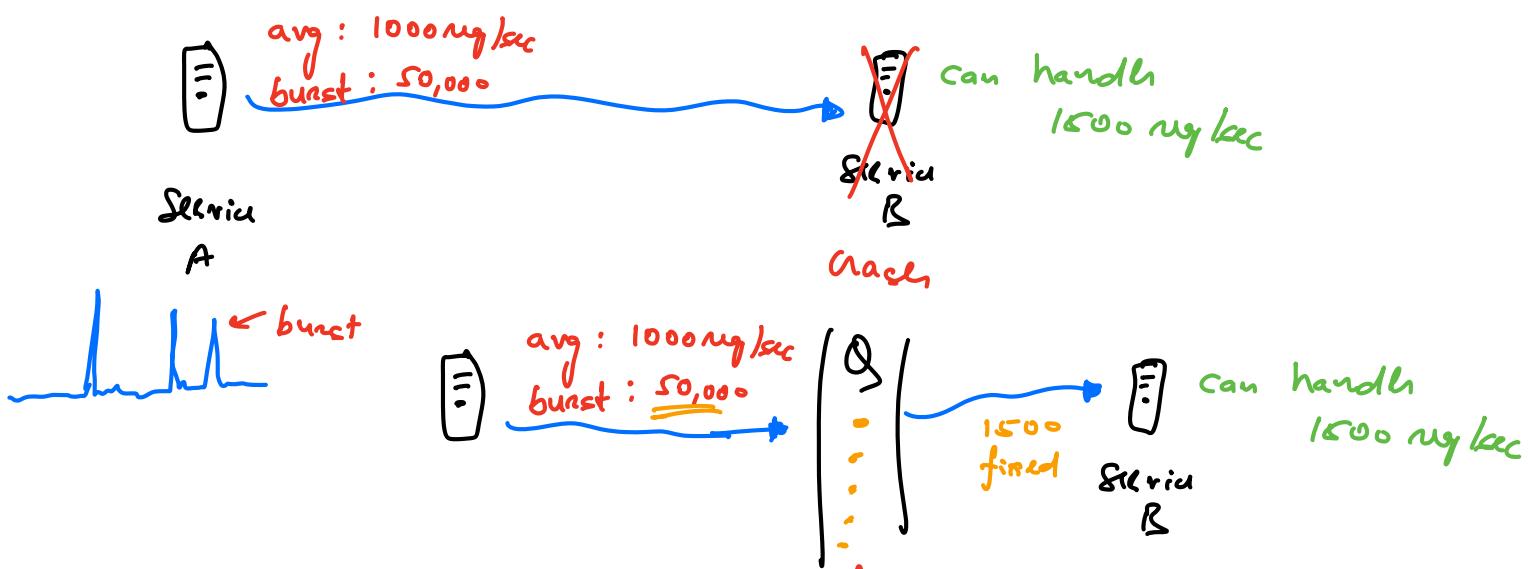
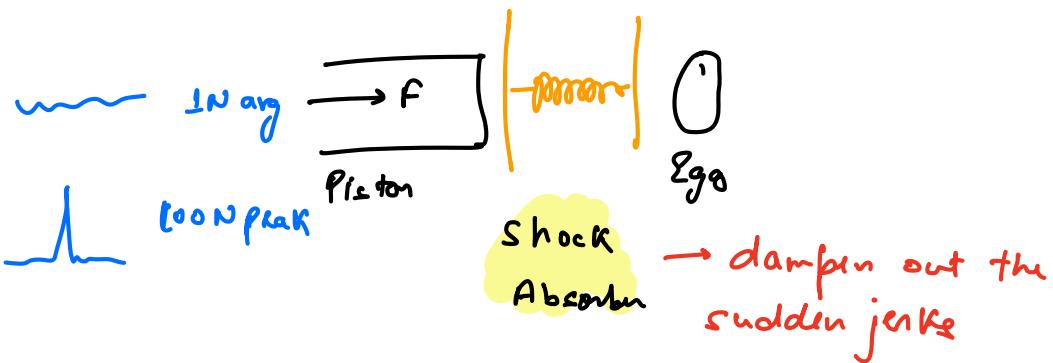
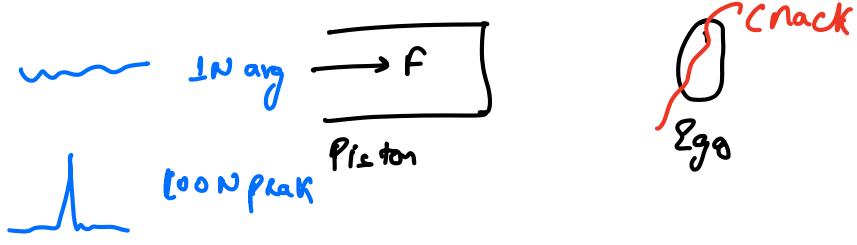
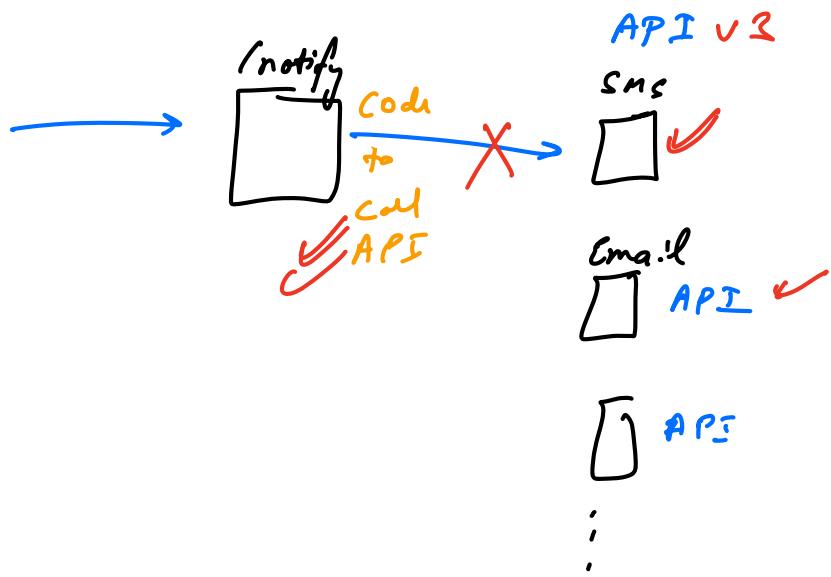


- ① user is online & connected → send notification ASAP
 - ② user can be offline → they will come back online after some time
 - any notifications that occurred when they were offline
- Only for in-app

Microservices → Event Driven Architecture
 Publisher Subscriber Model
 aka Producer Consumer

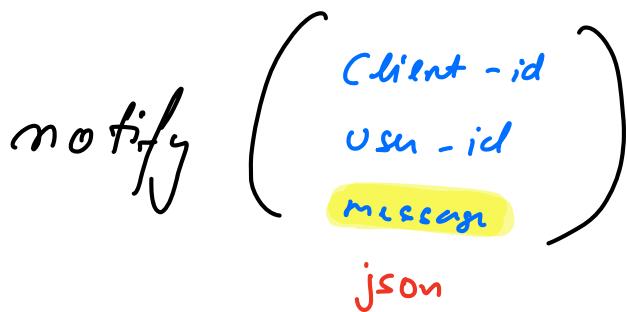
Pub sub = Producer Consumer
 ≠ Observers ≠ Event Handling

via
Message
Queue
like
Kafka



Q acts as a buffer

- ↳ prevents notification from getting dropped
- ↳ prevents service from crashing



msg :

type : OTP / Marketing / Finance / News
Used for priority setting

channels : {

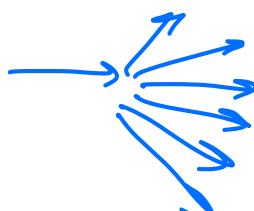
sms : { msg : — emoji : — }

email : { subject : ... body : ... }

} in-app : { msg : — icon : — }

action - url : 'Open(scren 2)' }

fan-out : 1 to many



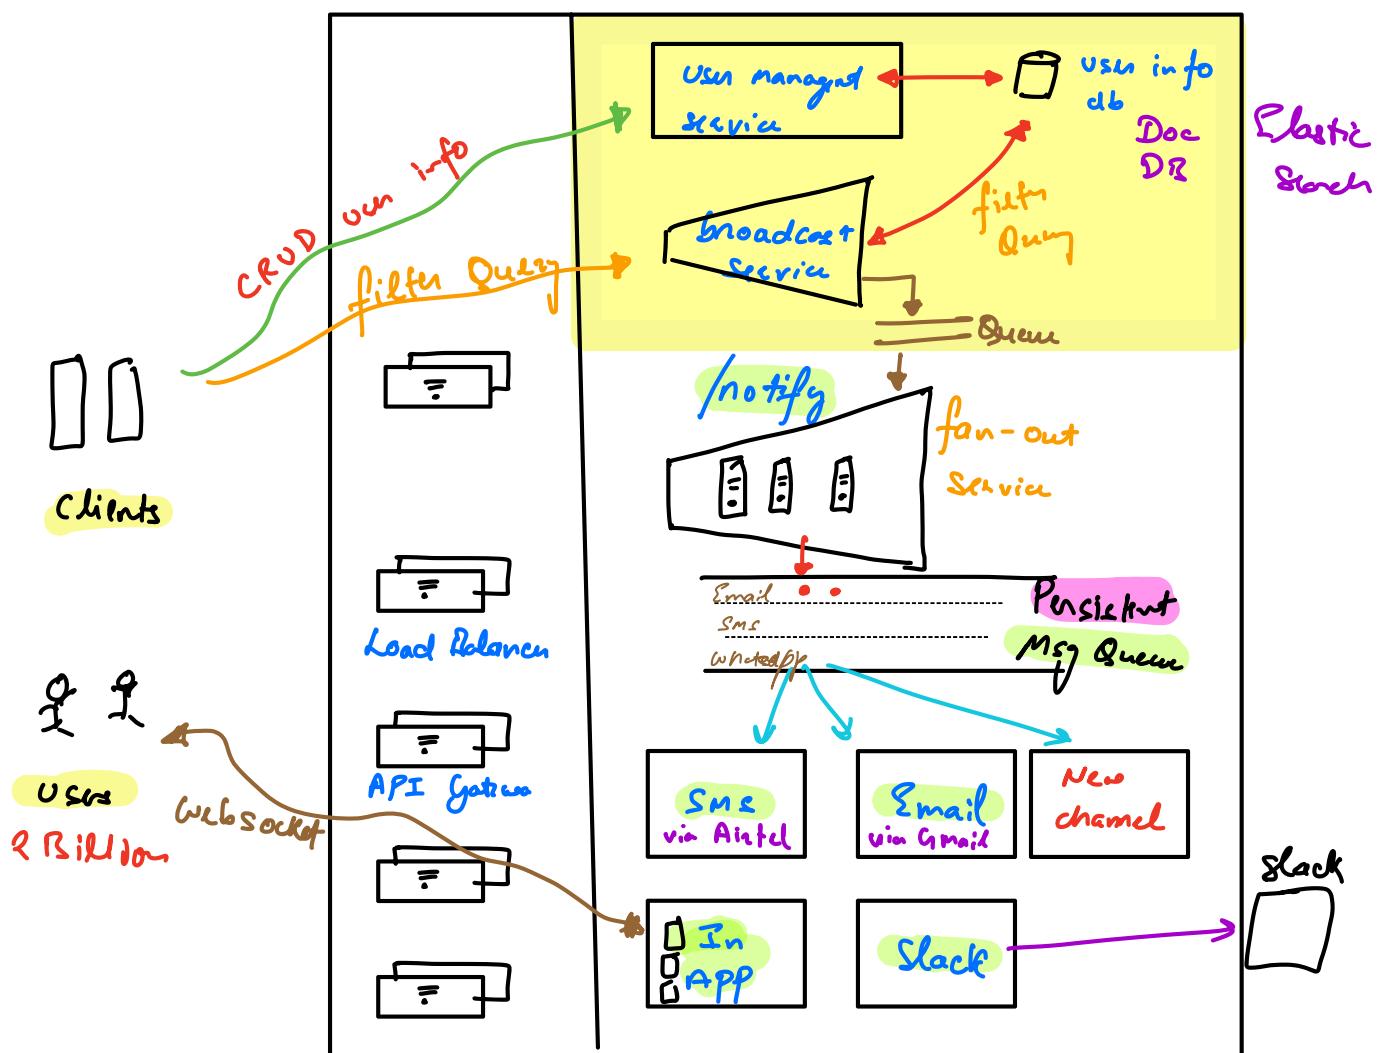
Partition the Msg Queue

→ handle load
ensure msg order

Create various topics — one for each channel
(Kafka)

each topic can be partitioned based on the avg load in that channel

② clients should be able to broadcast notification



choice of DB for client - user info

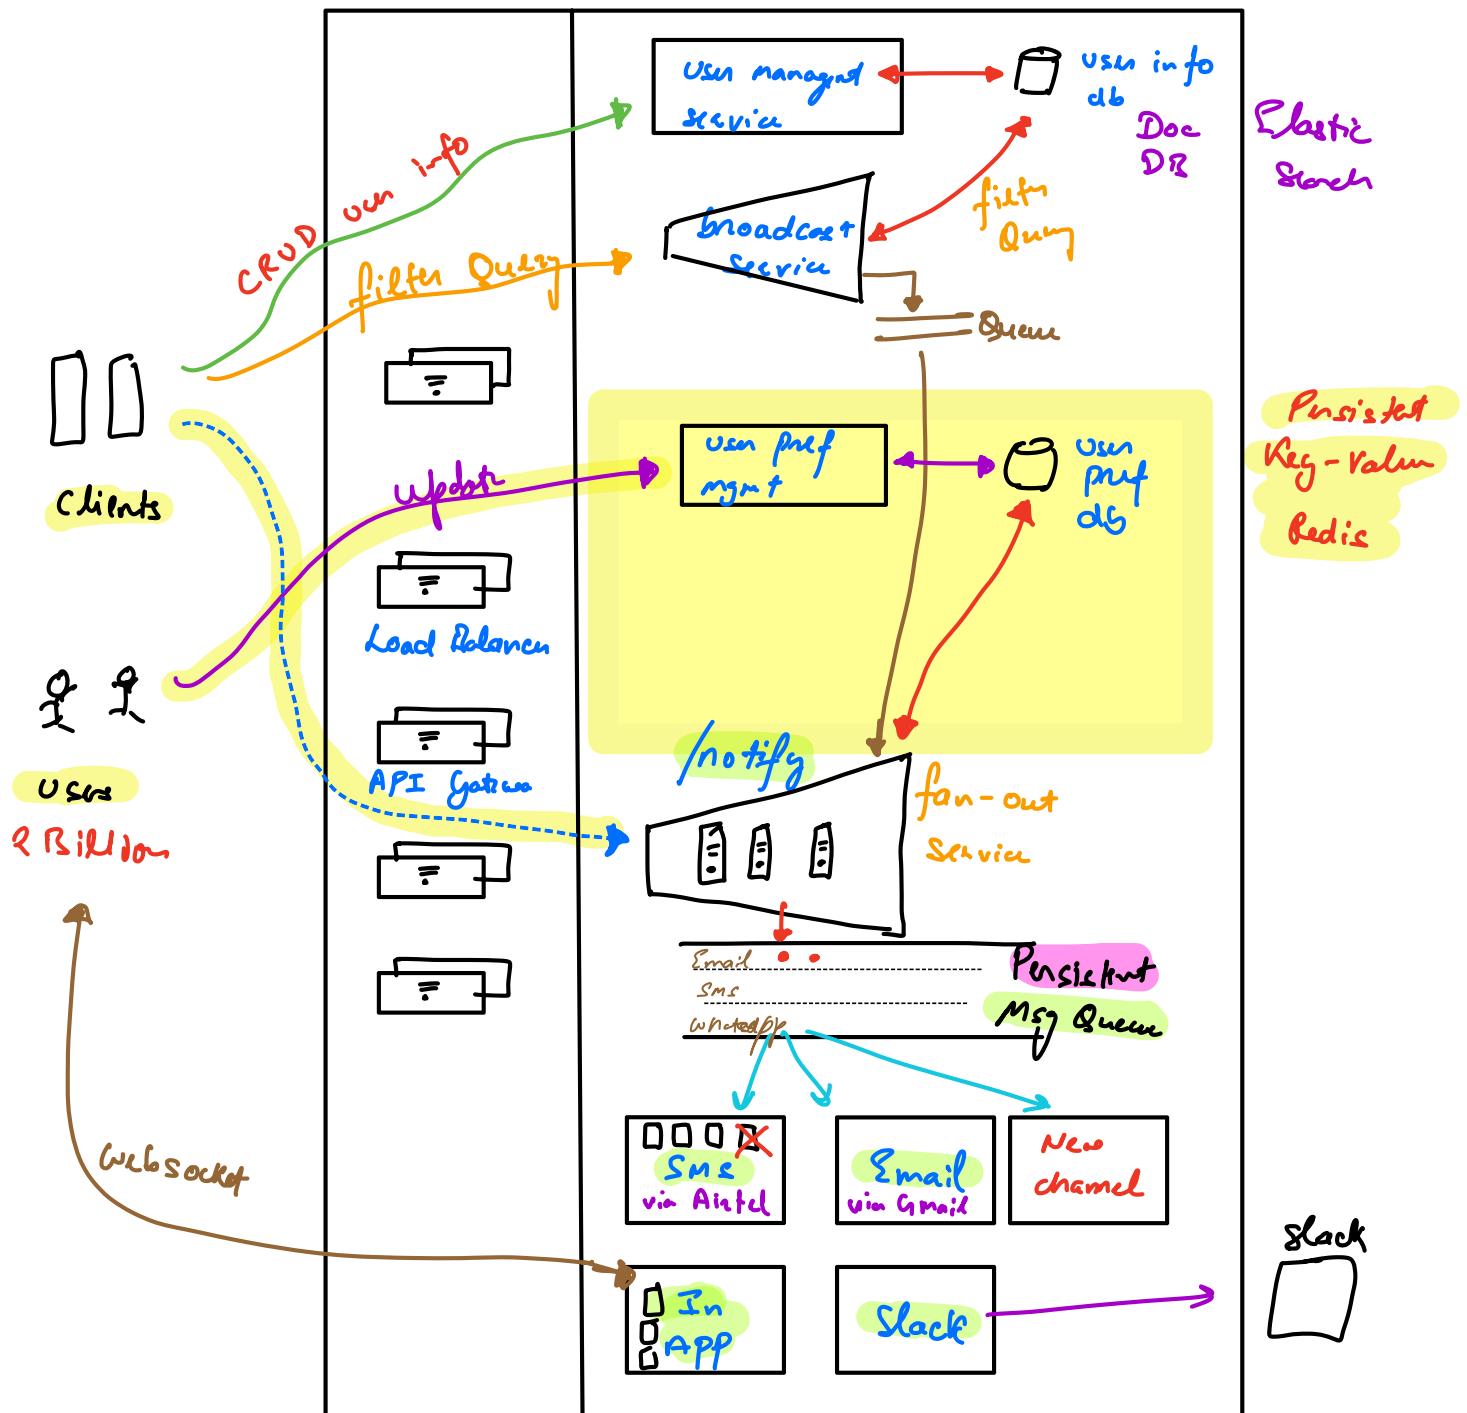
	Example	Pros	Cons
SQL	Postgres MySQL SQLite	<ul style="list-style-type: none"> • ACID • Normalization • joins • complex queries • relations • ... 	<ul style="list-style-type: none"> • Rigid schema • don't scale horizontally (unless you shard correctly)

NoSQL — built for scale (automatic sharding, replication...)

Document	Mongo DB Couch DB Cassandra DB Elastic Search	<ul style="list-style-type: none"> • Unstructured data • Full text search 	<ul style="list-style-type: none"> • no global index • no relations • simple queries
Column	Cassandra Scylla Big Table Hbase	<ul style="list-style-type: none"> • fast analytics (aggregation) • pagination by time • ultra high writes 	<ul style="list-style-type: none"> • no relations • simple queries • reading via multiple cols is slow
Key-Value	Redis Memcached Dynamo DB	<ul style="list-style-type: none"> • v. fast reads & writes • v. simple 	<ul style="list-style-type: none"> • no schema • no relations • (volatile)

③

Users should be able to manage notification preferences



choice of db for user pref data?

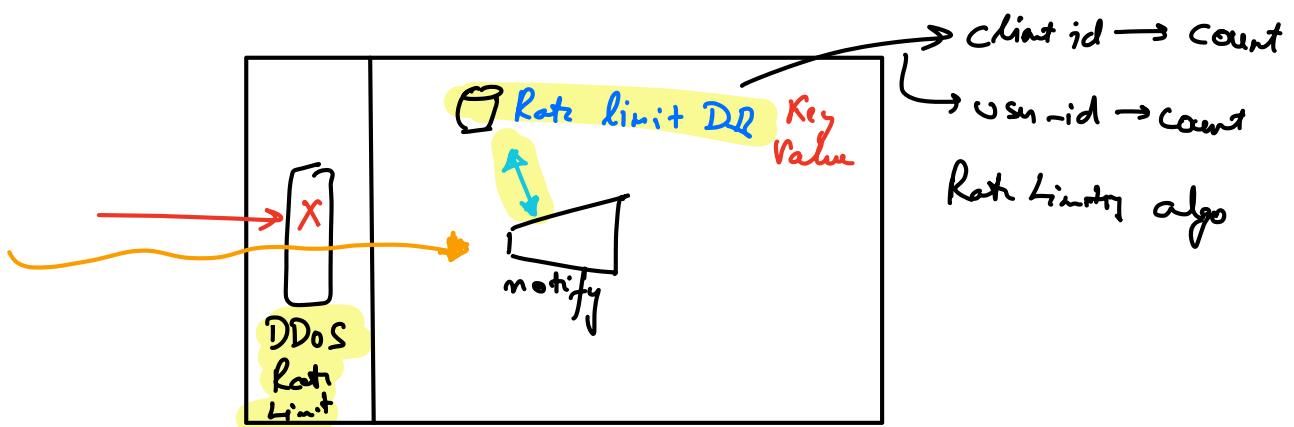
- Simple data — Given user whether a particular channel is enabled or not

User → channel → true/false

- 10M reads/sec

Key-Value (Redis + Persistence)

⑤ Rate Limiting





Observability

